

Title: Parallelization of Disease Propagation Algorithm (DPA)

Sungsoo Kim (EID: sk38224)

Abstract

The advent of COVID-19 novel virus from Wuhan, China and its spread across the world is making devastating humanitarian and economic impact. With its ever-increasing growth rate, it poses a large challenge to healthcare officials and others to gauge its effect and deploy countermeasures to control its spread. Such situation has inspired me to develop a disease propagation modeling algorithm to provide a better understanding of the COVID-19 virus. This algorithm will be based on C++ and OpenMP to maximize performance.

The project is divided into two major components: 1. Creating a general disease propagation model and 2. Parallelizing the code for optimal computation. The algorithm will be based on the SIR (Susceptible, Infectious and Recovered) model and other specific details of these sections will be determined through additional research. This research phase will help identify other major variables to consider and further enhance the accuracy of the simulator. Moreover, reference to existing mathematical models of disease propagation will also provide a general guidance in the development of the serial code.

Introduction

The simulation of the disease propagation model is developed based two major classes: Population and Person Additionally, there are system variables that govern the simulation. First, the system variables are set as the following:

System Variables

Field Name	Description	Value
INFECTION_RATE	Possibility of Infection when in contact	0.50
INTERACTION_NUM	Mean, S.D. of Interactions (per day)	Mean: 12, S.D: 3
DORMANT_PERIOD	Average Dormant Period of Virus	14 day
RECOVERY_PERIOD	Mean, S.D. of Recovery	Mean: 14, S.D.: 2
CHANCES	Random Chance of Infection	$0 < \text{CHANCES} < 1$ (Random)

Person Class

Field Name	Description	Value
Status	Current Status of the Patient	-
- Infected	The individual will be assumed not in self-isolation up to day 14. Once discovered, the individual will be recovering based on the recoveryDate field.	Status >= 1
- Recovered	Individual has recovered from the virus and gained immunity	Status = -1
- Susceptible	Individual has no immunity and is susceptible to the virus	Status = 0
- Inoculated	Individual has immunity to the virus	Status = -2
Interaction	Interaction count calculated based on national normal distribution. The numbers are renewed every day for accurate simulation	Individual Normal Distribution values based on the INTERACTION_NUM field
recoveryDate	Recovery date calculated based on the normal distribution for COVID-19. The value is initialized with the People Vector.	Individual Normal Distribution values based on the INTERACTION_NUM field

Population Class

Field Name	Description	Value
Size	Number of People in the Community	10,000,000
People <Vector>	Vector of Person objects with a size of "Size"	<Vector>
Day	Current Day of the Simulation	0 to user defined date

Output File

Name	Format
Disease_Simulation.csv	Day Total Sickness Uninfected Recovered & Inoculated

To optimize the simulation for the purpose of this project to parallelize, there were several hypotheses taken.

1. The modeling is based on a single community of people, and there are no external interactions with any interaction with the external communities.
2. Infected individuals will continue the social interaction during the virus dormant period (14 days for COVID-19) and then progress into the self-isolation period when he/she will begin the recovery.

3. There is no mutation to the virus – Once the individual has recovered from the virus, the individual is then immune
4. The chances of getting infected is kept constant throughout the simulation (Regardless of age, gender and other possible factors)
5. No external social act limitations are in place – No social distancing & Isolation except for the infected individuals with signs of infection.
6. Out of the interactions, we assume that 10% of them are “new” interactions whereas the rest are composed of our close friends, including our family, neighbors and others that we interact on a daily / bi-daily basis.

In terms of the variables, the population was set to 4,000,000 people to simulate relatively large sized city, the recovery and the dormant periods were taken from the WHO’s latest report on COVID-19 (April, 20’). Additionally, the chances of infections during an interaction with a virus holder was randomly selected to 30% and the number of interactions per day was taken from a study (Martire et al., 2018).

Project Objectives

1. *Code Development*
 - A. Create a Disease Propagation Algorithm based on the SIR model
 - B. Create a scalable solution that can vary the scope of the simulation
 - i. Introduce Objects *E.g. Nations, states etc.*
 - C. Compare output values to known data (*Johns Hopkins Coronavirus Research Center*)
2. *OpenMP Implementation*
 - A. Implement OpenMP functions to enable parallelization of the code
 - B. Generate a benchmark report to optimize parallel performance of the code
 - C.

Implementation

The simulation process can be divided into four different parts: Initialization of the People vector, setting initial conditions, virus simulation, and Counting. Here, the report will go through each part in detail and discuss how the program was converted from a serial to a parallel code.

1. Initialization of the People Vector

The program holds the information of a single person in a vector. A vector is a dynamic datatype and the initialization of the People vector took relatively less compared to the other hefty processes taken in the program. Therefore, the parallelization for this section was not done. However, in the case that the program is extended to be used for multiple communities (using multiple Populations class), it is possible to place a OpenMP for loop to iterate through different population classes to reduce the computation time.

2. Setting Initial Conditions

Once the People vector is constructed, the program takes in the initial conditions set by its user. The user specifies the number of infected people, number of inoculated people. Initially, the setup of the two conditions were done in two for loops. However, to optimize it for OpenMP application, the two for loops were combined.

```
// Initial Conditions
// =====
void setInitialCond(long inoculate_num, long infection_num) {
    printf("Setting Initial Conditions \n");
    long final_num = infection_num + inoculate_num;

    #ifdef _OPENMP
    #pragma omp parallel for schedule(runtime)
    for (long i = 0; i < final_num; i++) {
        if (i < infection_num) {
            people.at(i).setStatus(1);
        } else {
            people.at(i).setStatus(-2);
        }
    }
}
```

The initialization function takes in the two variables and goes through the for loop as the above. Because the two inputs are independent from each other, we can parallelize the loop without any race conditions. In the future, the program would be more effective if the initial conditions are set with the initialization stage.

3. Virus Simulation

The virus simulation is the section of the program that benefits the most from parallelization. This section divides the iteration of people in the population class and simulate human interaction person by person. Based on the variables and hypothesis outlined above, the simulate function runs the simulateDailyInteraction() function to go through an individual's interpersonal activity for the day.

```
#pragma omp parallel for schedule(dynamic, 1000)
for (long i = 0; i < people.size(); i++) {
    people[i].setInteraction(setInteractionNum()); // Adjusting Interaction
    simulateDailyInteraction(i); // Simulating Daily Interaction
}
```

As can be seen above, the OMP for loop is implemented to parallelize the operation. The index of the for loop indicate each person in the population. For this operation, there are chances of having a race condition. However, because most of the interactions are designed to be done to the neighboring people in the memory, a dynamic

scheduler is used. By using this method, the program will have less chance of having a race condition.

4. Counting Population

The last section of the program that implements OpenMP is the counting section. Here the program goes through the Population Vector and counts the status as it goes through each person. The implementation of OMP for loop was done as below.

```
#pragma omp parallel for schedule(dynamic) reduction (+:countR,countN,countP,countI)
for ( long i = 0; i < people.size(); i++) {
    status = people.at(i).getStatus();
    if ( status < 0 ) {
        countR++;
    } else if ( status == 0 ) {
        countN++;
    } else if ( status < DORMANT_PERIOD ) {
        countP++;
    } else {
        countI++;
    }
}
```

Unlike other for loops in the program, the counters in each thread need to be summed by the end of the execution. Hence, this for loop introduces the reduction for all the counters. By having this, all the counters in each thread can be added up to provide a total status of the population on a given date.

5. Others

Throughout the program, random number generation functions are used to calculate the possibility of infection and to find random encounters. For the original version of the code, it was implementing the rand() function, which was not thread safe and interfered with the performance of the parallelized program. As a result, the execution time became significantly larger than the serial execution time. Hence, these random functions were replaced by the rand_r function, which takes the seed from each thread to ensure that the threads are running the random functions independently and not conflict with each other. However, due to the way rand_r() functions are set, it tended to create the same values throughout its execution so it will require some tuning for the parallelization version to be used yet it would not influence the results for this experiment.

Experiment

The main objective of the experiment is to investigate the speed up of the program with different number of threads. Here are the general outlines of the controlled variables.

Independent Variable	Specifications
Number of Threads	1(Serial & Parallel), 2, 4, 8, 16, 32, 64, 128, 256

Dependent Variable	Specifications
Execution time	Time taken for Initial Conditions & Simulation

Controlled Variables	Specifications
Population	4,000,000 People
Number of Inoculated People	1000 People
Initial Infected Individuals	2 People
Simulation Length	35 days
Hardware	TACC Stampede2 skx-dev


Results

1. Serial vs. Parallel (1 Thread Execution)

	Serial Execution	Parallel Execution
Initial Conditions (s)	0.000398	0.000381
Simulation (s)	259.366	259.636


In the case of the single thread execution, the serial code and the parallelized code came out to have similar performance results. For the initial condition setup, the parallel execution resulted with better performance, whereas the serial execution of the code came to be slightly faster in the disease propagation simulation.

2. Parallel Performance in Initial Conditions




Interestingly, the time taken for initial conditions started increasing after 4 threads. While this can be an odd phenomenon to observe, we can deduce that the time required to process the section increases with more threads because there isn't enough tasks for the threads to process, making it rather inefficient to have more threads and to coordinate than having a few threads with less need to allocate tasks.

3. Parallel Performance in Simulation



Here we can see that the number of threads is inversely proportionate with the simulation execution time. Additionally, there were noticeable changes in the first few stages of the increments whereas there was relatively less impact on the execution speed as the experiment increased the number of threads.



If we take the inverse of the execution time, the relationship between the execution speed of the parallel program and the number of threads can be found. From this graph, it seems that there is a logarithmic relationship between the two variables. Yet, it evidently illustrates that the execution speed continues to rise as the number of threads increases.

Summary

As expected, the program displayed an inversely proportionate relationship for the simulation section of the software. On the other hand, the initial condition function showed that it actually becomes less effective to use more threads if the work shared is not big enough. Furthermore, during the conversion of the program, there were issues with the randomized number system where it could not generate truly unique set of random numbers. Hence, as of now, only the output generated from the serial version of the software creates can be considered useful. However, if this issue is resolved, the program will be able to handle even bigger populations and potentially start simulating interactions between different communities as well.

In terms of improvements, the program can combine the initialization of the Population vector with the initial condition setter to reduce the computational cost. To do this, the program will need to pre-allocate the memory for the vector. This way, each thread can independently access different elements of the vector and the program can fully benefit from parallelization.

References

Zhaoyang, R., Sliwinski, M., Martire, L., Smyth, J., & Zhaoyang, R. (2018). Age differences in adults' daily social interactions: An ecological momentary assessment study. *Psychology and Aging, 33*(4), 607–618. <https://doi.org/10.1037/pag0000242>

Appendix

Program Output

```
< SIMULATION >

Population: 4000000   Infection: 2   Inoculated: 1000
Setting Initial Conditions
< DAY 0 >   Inf: 2           Non-inf: 3998998   Self-isolation : 0       Rec OR Inn: 1000
< DAY 1 >   Inf: 4           Non-inf: 3998996   Self-isolation : 0       Rec OR Inn: 1000
< DAY 2 >   Inf: 20          Non-inf: 3998980   Self-isolation : 0       Rec OR Inn: 1000
< DAY 3 >   Inf: 115         Non-inf: 3998885   Self-isolation : 0       Rec OR Inn: 1000
< DAY 4 >   Inf: 539         Non-inf: 3998461   Self-isolation : 0       Rec OR Inn: 1000
< DAY 5 >   Inf: 2466        Non-inf: 3996534   Self-isolation : 0       Rec OR Inn: 1000
< DAY 6 >   Inf: 10966       Non-inf: 3988034   Self-isolation : 0       Rec OR Inn: 1000
< DAY 7 >   Inf: 49625       Non-inf: 3949375   Self-isolation : 0       Rec OR Inn: 1000
< DAY 8 >   Inf: 217109      Non-inf: 3781889   Self-isolation : 2       Rec OR Inn: 1000
< DAY 9 >   Inf: 831447      Non-inf: 3167548   Self-isolation : 5       Rec OR Inn: 1000
< DAY 10 >  Inf: 2219064     Non-inf: 1779903   Self-isolation : 33      Rec OR Inn: 1000
< DAY 11 >  Inf: 3501947     Non-inf: 496866    Self-isolation : 187     Rec OR Inn: 1000
< DAY 12 >  Inf: 3923619     Non-inf: 74507     Self-isolation : 874     Rec OR Inn: 1000
< DAY 13 >  Inf: 3986255     Non-inf: 8737      Self-isolation : 4008     Rec OR Inn: 1000
< DAY 14 >  Inf: 3980002     Non-inf: 903       Self-isolation : 18095    Rec OR Inn: 1000
< DAY 15 >  Inf: 3918086     Non-inf: 98        Self-isolation : 80816    Rec OR Inn: 1000
< DAY 16 >  Inf: 3656893     Non-inf: 6         Self-isolation : 342101   Rec OR Inn: 1000
< DAY 17 >  Inf: 2815083     Non-inf: 0         Self-isolation : 1183917  Rec OR Inn: 1000
< DAY 18 >  Inf: 1350231     Non-inf: 0         Self-isolation : 2648769  Rec OR Inn: 1000
< DAY 19 >  Inf: 326192      Non-inf: 0         Self-isolation : 3672808  Rec OR Inn: 1000
< DAY 20 >  Inf: 46188       Non-inf: 0         Self-isolation : 3952812  Rec OR Inn: 1000
< DAY 21 >  Inf: 5346        Non-inf: 0         Self-isolation : 3993654  Rec OR Inn: 1000
< DAY 22 >  Inf: 545         Non-inf: 0         Self-isolation : 3998455  Rec OR Inn: 1000
< DAY 23 >  Inf: 58          Non-inf: 0         Self-isolation : 3998940  Rec OR Inn: 1002
< DAY 24 >  Inf: 1           Non-inf: 0         Self-isolation : 3998994  Rec OR Inn: 1005
< DAY 25 >  Inf: 0           Non-inf: 0         Self-isolation : 3998967  Rec OR Inn: 1033
< DAY 26 >  Inf: 0           Non-inf: 0         Self-isolation : 3998813  Rec OR Inn: 1187
< DAY 27 >  Inf: 0           Non-inf: 0         Self-isolation : 3998126  Rec OR Inn: 1874
< DAY 28 >  Inf: 0           Non-inf: 0         Self-isolation : 3994992  Rec OR Inn: 5008
< DAY 29 >  Inf: 0           Non-inf: 0         Self-isolation : 3980905  Rec OR Inn: 19095
< DAY 30 >  Inf: 0           Non-inf: 0         Self-isolation : 3918184  Rec OR Inn: 81816
< DAY 31 >  Inf: 0           Non-inf: 0         Self-isolation : 3656899  Rec OR Inn: 343101
< DAY 32 >  Inf: 0           Non-inf: 0         Self-isolation : 2815083  Rec OR Inn: 1184917
< DAY 33 >  Inf: 0           Non-inf: 0         Self-isolation : 1350231  Rec OR Inn: 2649769
< DAY 34 >  Inf: 0           Non-inf: 0         Self-isolation : 326192   Rec OR Inn: 3673808
< DAY 35 >  Inf: 0           Non-inf: 0         Self-isolation : 46188    Rec OR Inn: 3953812

Data saved to file!
Population Generation :      0.0733058 seconds
Initial Conditions Set:     0.000175953 seconds
```